

Ames Housing Data

Thursday, 5 April 2018 5:09 PM

Regression and Classification with the Ames Housing Data

Github contents

https://github.com/prickofdeath/ames_housing_data/tree/master

	Filename	Description
1.	Project3_1	Data load, EDA
2.	Project3_2	Fixed Features Model using Linear Regression
3.	Project3_3	Renovatible Features Model using Linear Regression
4.	Project3_4	Identifying Abnorml sales with KNN
5.	Project3_5	Identifying Abnorml sales with regit regression

Introduction

You have just joined a new "full stack" real estate company in Ames, Iowa. The strategy of the firm is two-fold:

- Own the entire process from the purchase of the land all the way to sale of the house, and anything in between.
- Use statistical analysis to optimize investment and maximize return.

The company is still small, and though investment is substantial the short-term goals of the company are more oriented towards purchasing existing houses and flipping them as opposed to constructing entirely new houses. That being said, the company has access to a large construction workforce operating at rock-bottom prices.

1. Estimating the value of homes from fixed characteristics. Your superiors have outlined this year's strategy for the company:

Develop an algorithm to reliably estimate the value of residential houses based on fixed characteristics. Identify characteristics of houses that the company can cost-effectively change/renovate with their construction team. Evaluate the mean dollar value of different renovations. Then we can use that to buy houses that are likely to sell for more than the cost of the purchase plus renovations.

Your first job is to tackle #1. You have a dataset of housing sale data with a huge amount of features identifying different aspects of the house. The full description of the data features can be found in a separate file:

housing.csv data_description.txt You need to build a reliable estimator for the price of the house given characteristics of the house that cannot be renovated. Some examples include:

The neighborhood Square feet Bedrooms, bathrooms Basement and garage space and many more.

Some examples of things that ARE renovate-able:

Roof and exterior features "Quality" metrics, such as kitchen quality "Condition" metrics, such as condition of garage Heating and electrical components and generally anything you deem can be modified without having to undergo major construction on the house.

Goals:

1. Perform any cleaning, feature engineering, and EDA you deem necessary.
2. Be sure to remove any houses that are not residential from the dataset.
3. Identify fixed features that can predict price.
4. Train a model on pre-2010 data and evaluate its performance on the 2010 houses.
5. Characterize your model. How well does it perform? What are the best estimates of price?

Filter for only residential from mszoning

```
1 '''
2 MSZoning: Identifies the general zoning classification of the sale.
3
4 A—Agriculture
5 C—Commercial
6 FV—Floating Village Residential
7 I—Industrial
8 RH—Residential High Density
9 RL—Residential Low Density
10 RP—Residential Low Density Park
11 RM—Residential Medium Density
12
13 '''
14 # filter only residential from mszoning
15 commercial_zoning = ['A','C (all)','I']
16 house = house[~house['mszoning'].isin(commercial_zoning)]
17 house['mszoning'].value_counts() #check to ensure only residential remain
```



```
RL    1151
RM     218
FV      65
RH      16
Name: mszoning, dtype: int64
```

Check for null values

```
1 house.loc[:,pd.isnull(house).any()].head(5)
```


◀ ▶

Input Variable	fillna Decision	Rationale
lotfrontage	mean	
alley	drop	89 filled rows vs 1450 total
masvnrtype	drop	<pre>count 1442.000000 mean 104.404300 std 181.486539 min 0.000000 25% 0.000000 50% 0.000000 75% 166.750000 max 1600.000000 Name: masvnrarea, dtype: float64</pre> <p><matplotlib.axes._subplots.AxesSubplot at 0xecae278></p>  <p>The histogram displays the frequency distribution of the variable 'masvnrarea'. The x-axis represents the area values from 0 to 1600, and the y-axis represents the frequency from 0 to 1000. The distribution is highly right-skewed, with a very high frequency (over 1000) for values near zero, which rapidly decreases as the area increases, with a few small bars extending up to 1600.</p>
bsmtqual, bsmtcond, ... garage poolqc fence miscfeature	Convert null to 'no'	Blank means no basement
electrical	replace with SBrkr	Majority class
yearbuilt, remod_age	replace with age and remod_time,	

Part 2

Identify Price Drivers

Identify Price Drivers



Selecting Fixed Features

Predictor	Description
lotfrontage	linear feet of street connected to property
lotarea	lot size in sq ft
lotconfig	lot config eg Inside Lot, Corner Lot ...
neighborhood	physical location
bldgtype	type of dwelling
housestyle	style of dwelling
bsmtqual	height of basement
totalbsmtsf	total sq ft of basement area
1stflrsf	
2ndflrsf	
grlivarea	above grade ground living area sq ft
fullbath	
bedroomabvgr	
kitchenabvgr	
totrmsabvgrd	
fireplaces	
garagearea	garagecars removed, as garagearea and garagecars are highly correlated
wooddecksf	
openporchsf	
enclosedporchsf	
poolarea	
age	Calculated proxy for yearbuilt
remod_age	Calculated proxy for yearremodadd

Conversion of the non numeric features

- Where there is a rating (Ex, Gd, TA...), convert to 5 to 0 scale
- Otherwise pd.get_dummies

```
In [7]: 1 # Understanding the non numeric features that have been selected
2
3 print house_fixed['lotconfig'].unique(),'\n'
4 print house_fixed['neighborhood'].unique(),'\n'
5 print house_fixed['bldgtype'].unique(),'\n'
6 print house_fixed['housestyle'].unique(),'\n'
7 print house_fixed['bsmtqual'].unique()

['Inside' 'FR2' 'Corner' 'CuLDSac' 'FR3']

['CollgCr' 'Veenker' 'Crawfor' 'NoRidge' 'Mitchel' 'Somerst' 'NWAmes'
 'OldTown' 'BrkSide' 'Sawyer' 'NridgHt' 'Names' 'SawyerW' 'IDOTRR'
 'MeadowV' 'Edwards' 'Timber' 'Gilbert' 'StoneBr' 'ClearCr' 'NPKVill'
 'Blmngtn' 'BrDale' 'SWISU' 'Blueste']

['1Fam' '2fmCon' 'Duplex' 'TwnhsE' 'Twnhs']

['2Story' '1Story' '1.5Fin' '1.5Unf' 'SFoyer' 'SLvl' '2.5Unf' '2.5Fin']

['Gd' 'TA' 'Ex' 'No' 'Fa']
```

```
In [8]: 1 # map qualitative values to numeric values for bsmtqual
2 convert = {'Ex':5,'Gd':4,'TA':3,'Fa':2,'Po':1,'No':0}
3 house_fixed['bsmtqual'] = house['bsmtqual'].map(convert)
4
5
6 # Create df for the non numeric features
7 obj_df = house_fixed.select_dtypes(include=['object']).copy()
8 obj_df.head(4)
```

```
Out[8]:
```

	lotconfig	neighborhood	bldgtype	housestyle
0	Inside	CollgCr	1Fam	2Story
1	FR2	Veenker	1Fam	1Story
2	Inside	CollgCr	1Fam	2Story
3	Corner	Crawfor	1Fam	2Story

```
In [9]: 1 # binarize the non numeric features
2 dummy = pd.get_dummies(obj_df, columns=['lotconfig','neighborhood','bldgtype','housestyle'],
3         prefix=['lotconfig','neighbor','bldgtype','housestyle'], drop_first=True)
4 dummy.head(2)
```

```
Out[9]:
```

	lotconfig_CulDSac	lotconfig_FR2	lotconfig_FR3	lotconfig_Inside	neighbor_Blueste	neighbor_BrDale	neighbor_BrkSide	neighbor_ClearCr	neighbor_CollgCr
0	0	0	0	1	0	0	0	0	1
1	0	1	0	0	0	0	0	0	0

```
In [10]: 1 house_fixed = pd.concat([house_fixed,dummy],axis=1)
2 house_fixed.head(2)
3
```

```
Out[10]:
```

	lotfrontage	lotarea	lotconfig	neighborhood	bldgtype	housestyle	bsmtqual	totalbsmtsf	1stflrsf	2ndflrsf	grlivarea	fullbath	bedroomabvgr	kitchenabvgr	t
0	65.0	8450	Inside	CollgCr	1Fam	2Story	4	856	856	854	1710	2	3	1	
1	80.0	9600	FR2	Veenker	1Fam	1Story	4	1262	1262	0	1262	2	3	1	

Define test, train set by year (2010); Standardise prediction matrix

```
1 # define test and train sets
2
3 house_fixed_train = house_fixed[house_fixed.yrsold != 2010]
4 house_fixed_test = house_fixed[house_fixed.yrsold == 2010]
```

```
1 X_train = house_fixed_train[col_fixed_2].values
2 y_train = house_fixed_train['saleprice']
3
4 X_test = house_fixed_test[col_fixed_2].values
5 y_test = house_fixed_test['saleprice']
6
```

```
1 # standardise the test predictor matrix
2 from sklearn.preprocessing import StandardScaler
3 ss = StandardScaler()
4 Xs_train = ss.fit_transform(X_train)
5 print Xs_train.std(), Xs_train.mean()
6
7 # standardise the train predictor matrix
8 from sklearn.preprocessing import StandardScaler
9 ss = StandardScaler()
10 Xs_test = ss.fit_transform(X_test)
11 print Xs_test.std(), Xs_test.mean()
12
```


Run Linear Regression model

```
1 from sklearn import datasets, linear_model
2 from sklearn.metrics import mean_squared_error
3
4 lm = linear_model.LinearRegression()
5 model = lm.fit(Xs_train,y_train)
6
7 model_coef = model.coef_
8 predict = model.predict(Xs_train)
9 score = model.score(Xs_train,y_train)
10 RMSE_train = (mean_squared_error(y_train, predict))**0.5
11
12 print 'score', score, '\n'
13 print 'RMSE_train', RMSE_train
14
15 plt.figure(figsize=(15,8))
16 plt.scatter(predict, y_train)
17 plt.title('Model Prediction vs Actual', fontsize=20)
18 plt.xlabel ('Model Prediction')
19 plt.ylabel ('Actual')
20 plt.show()
```

score 0.82359878665

RMSE_train 33207.1023972

RMSLE_U1 0.111 33207.1023572



Check the Model Drivers

```
1 # checking the model drivers
2
3 model_coef = pd.DataFrame(zip(col_fixed_2,model.coef_),columns=['feature','coef']).sort_values(by=['coef'])
4 print 'Top 20 Model Drivers','\n','\n'
5 print model_coef.tail(10),'\n'
6 print model_coef.head(10)
7
8 #model_coef.sort_values(by=['coef'])
```

Top 20 Model Drivers

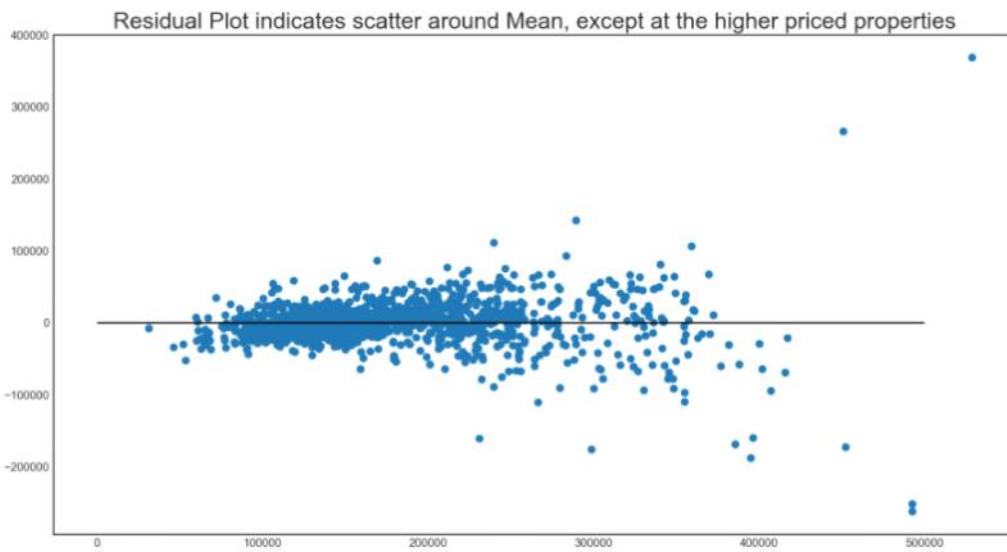
	feature	coef
10	totrmsabvgrd	5922.839012
12	garagearea	6640.852661
4	1stflrsf	7090.827599
45	neighbor_StoneBr	8958.390299
2	bsmtqual	9778.404982
53	housestyle_1Story	9874.225156
38	neighbor_NoRidge	10120.990207
5	2ndflrsf	12971.737896
39	neighbor_NridgHt	14342.750956
6	grlivarea	22861.668704

	feature	coef
51	bldgtype_TwnhsE	-8605.554716
18	age	-8130.602792
50	bldgtype_Twnhs	-7616.679995
19	remod_age	-7110.883346
8	bedroomabvgr	-6692.027508
9	kitchenabvgr	-6528.524257
30	neighbor_Edwards	-4514.974442
40	neighbor_OldTown	-3915.526239
37	neighbor_NWAmes	-3389.365781
31	neighbor_Gilbert	-3364.008959



Price Predictors

- Features that matter are the neighborhood, area, age, remod age

Check the Residual Plot to ensure no bias



Model vs Predict



Part 3 Determine Impact of Renovatable Features

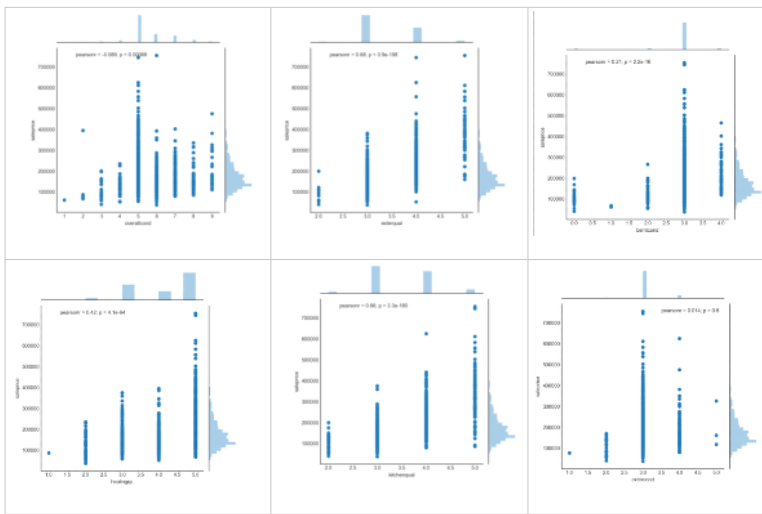
Selecting Renovatable Features

Renovatable Predictor	Description
overallcond	the overall condition of the house 10: very excellent, 1: very poor
exterqual	quality of material on the exterior
extercond	present condition of the ext material
bsmtcond	general condition of basement
heatingqc	heating quality and condition
kitchenqual	kitchen quality

Not including these as renovatable predictors as they also have a binary condition available or none and therefore not comparable

Renovatable Predictor	Description
fence	fence quality
poolqc	pool quality
garagecond	garage condition
fireplacequ	fireplace quality

Reviewing the jointplots



Binarize non numeric features

```

1 # binarize the non numeric features
2 dummy = pd.get_dummies(obj_df, columns=col,
3                       prefix = col, drop_first=True)
4 dummy.head(2)
5

```



Run the Model for the test-train

```

1 lm = linear_model.LinearRegression()
2 model = lm.fit(Xs_train,y_train)
3
4 predict = model.predict(Xs_train)
5 score = model.score(Xs_train,y_train)
6 print score
7
8 plt.figure(figsize=(15,10))
9 plt.scatter(predict, y_train)
10 plt.title('Model Prediction vs Actual', fontsize=20)
11 plt.xlabel ('Model Prediction')
12 plt.ylabel ('Actual')
13 plt.show()


```

0.919029045893




Running the statsmodel

Model Prediction vs Actual



Residual Plot indicates scatter around Mean, except at the higher priced properties



The Model now has a score of 91% vs 82% for fixed



Commentary

- Model without the renovatable features have a score of **0.82**
- Model including renovatable features have a score of **0.91**, indicating that renovatable features drive upside to the salesprice
- Objective: Look for undervalued properties, condition < TA
 - Predicted fixed_salesprice > salesprice_actual
 - Predicted fixed+renovatable_salesprice > salesprice_actual

Part 3a What Property characteristics predict an 'abnormal' sale?

The SaleCondition feature indicates the circumstances of the house sale. From the data file, we can see that the possibilities are:

Normal Normal Sale
 Abnorml Abnormal Sale - trade, foreclosure, short sale
 AdjLand Adjoining Land Purchase
 Alloca Allocation - two linked properties with separate deeds, typically condo with a garage unit
 Family Sale between family members
 Partial Home was not completed when last assessed (associated with New Homes)

One of the executives at your company has an "in" with higher-ups at the major regional bank. His friends at the bank have made him a proposal: if he can reliably indicate what features, if any, predict "abnormal" sales (foreclosures, short sales, etc.), then in return the bank will give him first dibs on the pre-auction purchase of those properties (at a dirt-cheap price). He has tasked you with determining (and adequately validating) which features of a property predict this type of sale.

Start with a simple KNN model

Encode target class variable salecondition Abnorml: 1, not Abnorml: 0

```
: 1 # encode the target class variable 'salecondition' to be a binary 1 if Abnorml, 0 if not Abnorml
  2 house['salecondition'] = house['salecondition'].map(lambda x: 1 if (x == 'Abnorml') else 0)
  3 house['salecondition'].value_counts()

: 0    1354
  1     96
  Name: salecondition, dtype: int64
```

Choose a subset of house that have a potential impact on salecondition

```

: 1 # choose a subset of house that have a potential impact on salecondition
2 col = ['age', 'remod_age', 'mssubclass', 'mszoning', 'lotarea', 'neighborhood', 'housestyle', 'salecondition']
3 h2 = house[col] #h2 is the house subset
4 h2.head(4)

```

```

:  age  remod_age  mssubclass  mszoning  lotarea  neighborhood  housestyle  salecondition
0    15         15         60      RL    8450      CollgCr      2Story         0
1    42         42         20      RL    9600      Veenker      1Story         0
2    17         16         60      RL   11250      CollgCr      2Story         0
3   103         48         70      RL    9550      Crawfor      2Story         1

```


Examine correlation structure of the dataset

```

1 # examine correlation structure of the dataset
2
3 fig,ax=plt.subplots(figsize=(25,10))
4 cmap=sns.diverging_palette(220,10, as_cmap=True)
5
6 h2_corr=h2.corr()
7 mask=np.zeros_like(h2_corr,dtype=np.bool)
8 mask[np.triu_indices_from(mask)] = True
9
10 plot_a = sns.heatmap(h2_corr, mask = mask, ax=ax, cmap=cmap)
11 plot_a.tick_params(labelsize = 16)
12 plot_a.set_title('Sale Condition Features', fontsize=30)

```

Text(0.5,1,u'Sale Condition Features')



Not seeing any correlation with the selected predictors

Baseline 0.93

```

1 # baseline accuracy
2 baseline = 1 - np.mean(y)
3 print 'baseline:', baseline

```

baseline: 0.933793103448

KNN with k = 5, 2, 10

Cross validation using StratifiedKFold (class imbalance)

Standardised, non-Standardised

Non data driven feature selection

No class balancing

Hypothesis#1: mssubclass, neighborhood, housestyle

Hypothesis#2: neighborhood only

Conclusion: KNN model does not provide a better outcome than baseline

Best + 0.0002% improvement on baseline

```

1 # hypothesis: neighborhood is factor
2 y = h2['salecondition']
3 X = h2[['neighborhood_Blueste', 'neighborhood_BrDale', 'neighborhood_BrkSide', 'neighborhood_Clear
4         'neighborhood_CollgCr', 'neighborhood_Crawfor', 'neighborhood_Edwards', 'neighborhood_Gilb
5         'neighborhood_IDOTRR', 'neighborhood_MeadowV', 'neighborhood_Mitchel', 'neighborhood_NAMES
6         'neighborhood_NPkVill', 'neighborhood_NWAmes', 'neighborhood_NoRidge', 'neighborhood_Nridg
7         'neighborhood_OldTown', 'neighborhood_SWISU', 'neighborhood_Sawyer', 'neighborhood_SawyerW
8         'neighborhood_Somerst', 'neighborhood_StoneBr', 'neighborhood_Timber', 'neighborhood_Veenk
9         'neighborhood_Blueste',
10        'neighborhood_BrDale', 'neighborhood_BrkSide', 'neighborhood_ClearCr', 'neighborhood_Collg
11        'neighborhood_Crawfor', 'neighborhood_Edwards', 'neighborhood_Gilbert', 'neighborhood_IDOT
12        'neighborhood_MeadowV', 'neighborhood_Mitchel', 'neighborhood_NAMES', 'neighborhood_NPkVil
13        'neighborhood_NWAmes', 'neighborhood_NoRidge', 'neighborhood_NridgHt', 'neighborhood_OldTo
14        'neighborhood_SWISU', 'neighborhood_Sawyer', 'neighborhood_SawyerW', 'neighborhood_Somerst
15        'neighborhood_StoneBr', 'neighborhood_Timber', 'neighborhood_Veenker', 'neighborhood_Blues
16        'neighborhood_BrkSide', 'neighborhood_ClearCr', 'neighborhood_CollgCr', 'neighborhood_Craw
17        'neighborhood_Edwards', 'neighborhood_Gilbert', 'neighborhood_IDOTRR', 'neighborhood_Mead
18        'neighborhood_Mitchel', 'neighborhood_NAMES', 'neighborhood_NPkVill', 'neighborhood_NWAmes
19        'neighborhood_NoRidge', 'neighborhood_NridgHt', 'neighborhood_OldTown', 'neighborhood_SWIS
20        'neighborhood_Sawyer', 'neighborhood_SawyerW', 'neighborhood_Somerst', 'neighborhood_Stone
21        'neighborhood_Timber', 'neighborhood_Veenker']]
22
23 # standardize the predictor matrix
24 from sklearn.preprocessing import StandardScaler
25 ss = StandardScaler()
26 Xs = ss.fit_transform(X)
27

```

```

1 # cross validate mean accuracy for a KNN with 5 neighbors
2 knn5 = KNeighborsClassifier(n_neighbors=5, weights='uniform')
3 scores = accuracy_crossvalidator(Xs, y, knn5, cv_indices)
4
5 print 'baseline:', baseline
6 percent_gain = (((np.mean(scores))/baseline)-1)*100
7 print 'Percent Gain: ', percent_gain, '%'

```

```

('Fold accuracy:', 0.93127147766323026)
('Fold accuracy:', 0.93103448275862066)
('Fold accuracy:', 0.93103448275862066)
('Fold accuracy:', 0.92758620689655169)
('Fold accuracy:', 0.9307958477508651)
('Mean CV accuracy:', 0.93034449956557774)
baseline: 0.933793103448
Percent Gain: -0.369311346375 %

```

```

1 # cross validate mean accuracy for a KNN with 2 neighbors
2 knn2 = KNeighborsClassifier(n_neighbors=2, weights='uniform')
3 scores = accuracy_crossvalidator(Xs, y, knn2, cv_indices)
4
5 print 'baseline:', baseline
6 percent_gain = (((np.mean(scores))/baseline)-1)*100
7 print 'Percent Gain: ', percent_gain, '%'

```

```

('Fold accuracy:', 0.92439862542955331)
('Fold accuracy:', 0.93103448275862066)
('Fold accuracy:', 0.93103448275862066)
('Fold accuracy:', 0.93103448275862066)
('Fold accuracy:', 0.92387543252595161)
('Mean CV accuracy:', 0.92827550124627334)
baseline: 0.933793103448
Percent Gain: -0.590880590318 %

```

```

1 # cross validate mean accuracy for a KNN with 10 neighbors
2 knn10 = KNeighborsClassifier(n_neighbors=10, weights='uniform')
3 scores = accuracy_crossvalidator(Xs, y, knn10, cv_indices)
4
5 print 'baseline:', baseline
6 percent_gain = (((np.mean(scores))/baseline)-1)*100
7 print 'Percent Gain: ', percent_gain, '%'

```

```

('Fold accuracy:', 0.93127147766323026)
('Fold accuracy:', 0.93448275862068964)
('Fold accuracy:', 0.93448275862068964)
('Fold accuracy:', 0.93448275862068964)
('Fold accuracy:', 0.93425605536332179)
('Mean CV accuracy:', 0.93379516177772415)
baseline: 0.933793103448
Percent Gain: 0.000220426713438 %

```

Part 3b What Property characteristics predict an 'abnormal' sale?

- A different approach to identifying predictors
- Logit regression
- Re-Run KNN

Use `groupby` `salecondition` to find meaningful mean differences that can indicate predictors

```

1 # Using groupby 'salecondition' to find meaningful differences that can indicate predictors
2 house.groupby('salecondition').mean()

```

	id	mssubclass	lotfrontage	lotarea	overallqual	overallcond	yearbuilt	yearremodadd	bsmtfins
salecondition									
0	730.500000	56.783604	70.177844	10595.706056	6.144756	5.591581	1972.340473	1985.725258	447.1329
1	743.552083	58.593750	68.290617	9510.104167	5.687500	5.447917	1961.052083	1975.520833	417.3750

Observations for Abnorml

In project3_4, the hypothesis was that age, remod_age, mssubclass, mszoning, lotarea, neighborhood, housestyle determine abnormal sales. Using KNN, it did not appear to be so.

Taking a different approach above to shortlist predictors for salecondition 1 = Abnorml, vs 0 = Not Abnorml, based on the groupby above

- lotarea is smaller
- overallqual, overallcond is lower
- yearbuilt, yearremodadd is older -> age, remod_age is higher
- lowqualfinsf is higher
- poolarea is higher
- saleprice is lower

Based on above, chose new features that may have a potential impact on salecondition


Examine correlations

```

1 # examine correlation structure of the dataset
2
3 fig,ax=plt.subplots(figsize=(25,10))
4 cmap=sns.diverging_palette(220,10, as_cmap=True)
5
6 h2_corr=h2.corr()
7 mask=np.zeros_like(h2_corr,dtype=np.bool)
8 mask[np.triu_indices_from(mask)] = True
9
10 plot_a = sns.heatmap(h2_corr, mask = mask, ax=ax, cmap=cmap)
11 plot_a.tick_params(labelsize = 12)
12 plot_a.set_title('Sale Condition Features', fontsize=30)

```

Text(0.5,1,u'Sale Condition Features')



Upsample minority class to balance

```

1 # up-sample minority class method to balance; resample with replacement
2 from sklearn.utils import resample
3
4 # seperate majority and minority classes
5 h2_majority = h2[h2.salecondition == 0]
6 h2_minority = h2[h2.salecondition == 1]
7
8 # upsample minority class
9 h2_minority_upsampled = resample(h2_minority, replace = True, n_samples = 1354, random_state=0)
10
11 # combine majority class with upsampled minority class
12 h2_upsampled = pd.concat([h2_majority, h2_minority_upsampled])
13
14 # check new class counts
15 h2_upsampled['salecondition'].value_counts()

```

```

1 1354
0 1354
Name: salecondition, dtype: int64

```

Run logit regression


```

1 import statsmodels.api as sm
2 logit_model=sm.Logit(y,X)
3 model = logit_model.fit()
4 print (model.summary())

```

Optimization terminated successfully.
 Current function value: 0.648121
 Iterations 6

Logit Regression Results

```

=====
Dep. Variable:          salecondition    No. Observations:          2708
Model:                  Logit           Df Residuals:              2700
Method:                 MLE             Df Model:                  7
Date:                  Tue, 03 Apr 2018   Pseudo R-squ.:            0.06496
Time:                  19:25:59          Log-Likelihood:           -1755.1
converged:              True             LL-Null:                  -1877.0
                                   LLR p-value:            5.604e-49
=====

```

	coef	std err	z	P> z	[0.025	0.975]
lotarea	-1.331e-05	7.93e-06	-1.679	0.093	-2.88e-05	2.23e-06
overallqual	0.1615	0.041	3.891	0.000	0.080	0.243
overallcond	-0.0913	0.035	-2.629	0.009	-0.159	-0.023
age	0.0033	0.002	1.630	0.103	-0.001	0.007
remod_age	0.0139	0.002	5.838	0.000	0.009	0.019
lowqualfinsf	0.0015	0.001	1.920	0.055	-3.14e-05	0.003
poolarea	0.0078	0.001	6.301	0.000	0.005	0.010
saleprice	-6.3e-06	1.1e-06	-5.736	0.000	-8.45e-06	-4.15e-06

- If p value < 0.05, consider the feature significant
- Based on z and corresponding p values, overallqual, overallcond, remod_age, poolarea are significant to model
- Drop lotarea, age, lowqualfinsf

Standardise predictor matrix
 Train_test_split
 Cross validation

```

1 # Cross Validation
2 # http://scikit-learn.org/stable/modules/classes.html#module-sklearn.model_selection
3
4 from sklearn import model_selection
5 from sklearn.model_selection import cross_val_score
6
7 # choice of stratified as the number of Abnorml sale condition is small
8 Skfold = model_selection.StratifiedKFold(n_splits=10, random_state=7)
9 model_cv = LogisticRegression()
10 results = model_selection.cross_val_score(model_cv, X_train, y_train, cv=Skfold)
11 print results
12 print 'average accuracy: {:.3f}'.format(results.mean())
13

```

```

[ 0.58115183  0.60209424  0.62105263  0.56084656  0.62962963  0.60846561
 0.65608466  0.6031746  0.65608466  0.61904762]
average accuracy: 0.614

```

Check model capability

- Confusion matrix
- Classification_report
- ROC, AUC

```

1 # confusion matrix
2 from sklearn.metrics import confusion_matrix
3 con = confusion_matrix(y_test, y_pred)
4
5 confusion = pd.DataFrame(con,index=['is not Abnorml','is Abnorml'], columns=['predicted not Abnorml','predicted Abnorml'])
6 confusion
7

```

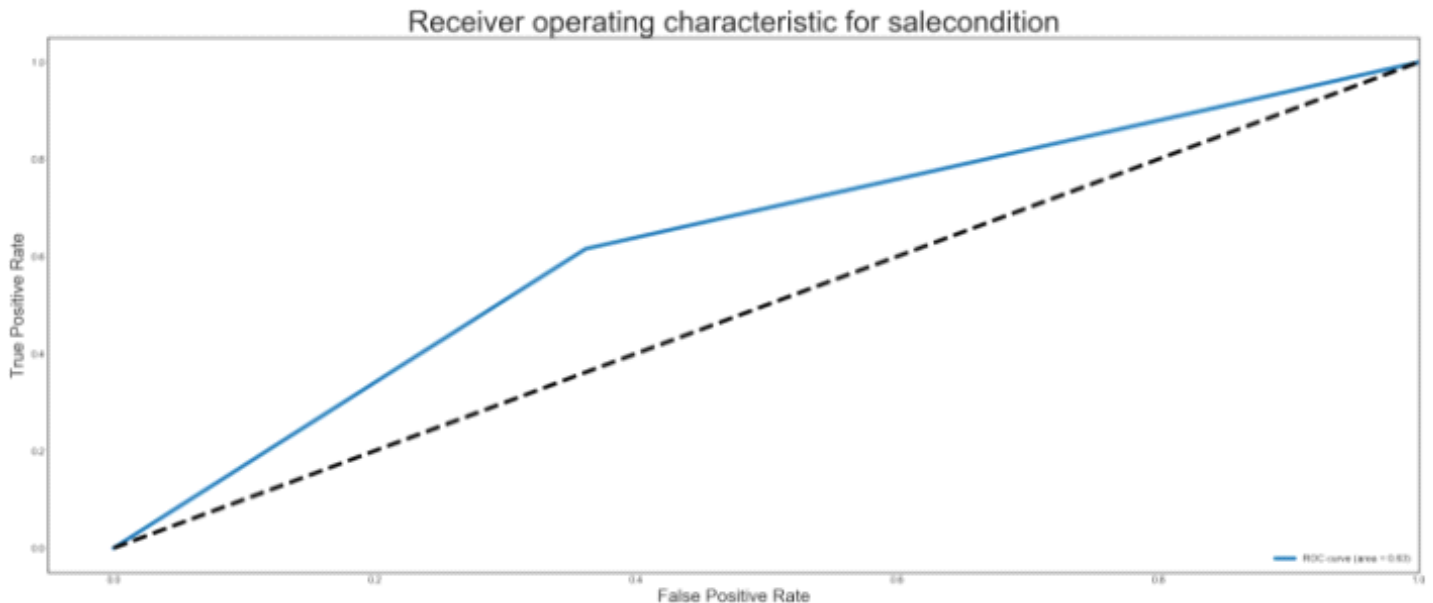
	predicted not Abnorml	predicted Abnorml
is not Abnorml	263	149
is Abnorml	154	247

```

1 from sklearn.metrics import classification_report
2 # precision - percent that model was correct when it was predicting true
3 # recall - percent that the model correctly predicted 1 when label was 1
4 # f1 score best = 1 worst = 0
5
6 print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.63	0.64	0.63	412
1	0.62	0.62	0.62	401
avg / total	0.63	0.63	0.63	813



Run the new feature set with KNN to see if it is better than the previous KNN model

StratifiedKFold, n_splits = 10
KNN, n_neighbors = 5, 2, 10

```

1 ## Instantiate the model with 5 neighbors.
2 knn = KNeighborsClassifier(n_neighbors=2)
3 ## Fit the model on the training data.
4 knn.fit(X_train, y_train)
5 ## See how the model performs on the test data.
6 print 'knn score:', knn.score(X_test, y_test), '\n'
7
8 Skfold = model_selection.StratifiedKFold(n_splits=10, random_state=7)
9 model_cv = KNeighborsClassifier(n_neighbors=2)
10 results = model_selection.cross_val_score(model_cv, X_train, y_train, cv=Skfold)
11 print results
12 print 'average accuracy: {:.3f}'.format(results.mean())

```

knn score: 0.960639606396

```

[ 0.97382199  0.92146597  0.96315789  0.95238095  0.94179894  0.94708995
  0.96825397  0.95767196  0.93650794  0.94179894]
average accuracy: 0.950

```

Revisiting the feature selection with RFE


```

1 from sklearn.feature_selection import RFE
2
3 logreg2 = LogisticRegression()
4 rfe = RFE(logreg,100)
5 rfe = rfe.fit(house[X],house[y])
6
7 play = sorted(zip(rfe.ranking_,X))
8 print play

```

```

[(1, 'bldgtype_Duplex'), (1, 'bldgtype_Twnhs'), (1, 'bsmtcond_Gd'), (1, 'bsmtcond_Po'), (1, 'bsmtexposure_Gd'), (1, 'bsmtexposure_Mn'), (1, 'bsmtfintype1_BLQ'), (1, 'bsmtfintype1_LwQ'), (1, 'bsmtfintype2_BLQ'), (1, 'bsmtfintype2_LwQ'), (1, 'bsmthalfbath'), (1, 'bsmtqual_Fa'), (1, 'bsmtqual_TA'), (1, 'centralair_Y'), (1, 'condition1_Feetr'), (1, 'condition1_PosN'), (1, 'condition1_RRAe'), (1, 'electrical_Mix'), (1, 'electrical_SBrkr'), (1, 'extercond_Fa'), (1, 'extercond_Gd'), (1, 'extercond_TA'), (1, 'exterior1st_BrkComm'), (1, 'exterior1st_CemntBd'), (1, 'exterior1st_Stone'), (1, 'exterior1st_Stucco'), (1, 'exterior1st_VinylSd'), (1, 'exterior2nd_BrkFace'), (1, 'exterior2nd_HdBoard'), (1, 'exterior2nd_IrStucc'), (1, 'exterior2nd_Stone'), (1, 'exterior2nd_Wd Sdng'), (1, 'exterior2nd_Wd Shng'), (1, 'fence_GdWo'), (1, 'fence_MnPrv'), (1, 'fence_No'), (1, 'fireplacequ_Fa'), (1, 'fireplacequ_Gd'), (1, 'fireplacequ_Po'), (1, 'fireplaces'), (1, 'foundation_Slab'), (1, 'fullbath'), (1, 'functional_Min2'), (1, 'functional_Sev'), (1, 'garagecond_Fa'), (1, 'garagecond_Po'), (1, 'garagefinish_No'), (1, 'garagequal_Fa'), (1, 'garagequal_Gd'), (1, 'garagequal_Po'), (1, 'garagequal_TA'), (1, 'garagetype_Attchd'), (1, 'garagetype_Basment'), (1, 'garagetype_CarPort'), (1, 'garagetype_Detachd'), (1, 'heating_GasA'), (1, 'heating_GasW'), (1, 'heatingqc_Fa'), (1, 'heatingqc_Gd'), (1, 'housestyle_1.5Unf'), (1, 'housestyle_2.5Fin'), (1, 'housestyle_2.5Unf'), (1, 'housestyle_SLvl'), (1, 'kitchenqual_Fa'), (1, 'kitchenqual_TA'), (1, 'landcontour_HLS'), (1, 'landcontour_Low'), (1, 'landslope_Mod'), (1, 'landslope_Sev'), (1, 'lotconfig_FR2'), (1, 'lotconfig_Inside'), (1, 'miscfeature_No'), (1, 'mszoning_RH'), (1, 'mszoning_RL'), (1, 'mszoning_RM'), (1, 'neighborhood_BrDale'), (1, 'neighborhood_ClearCr'), (1, 'neighborhood_CollCr'), (1, 'neighborhood_Edwards'), (1, 'neighborhood_Gilbert'), (1, 'neighborhood_IDOTRR'), (1, 'neighborhood_NAmes'), (1, 'neighborhood_NWAmes'), (1, 'neighborhood_NoRidge'), (1, 'neighborhood_NridgHt'), (1, 'neighborhood_Somerst'), (1, 'neighborhood_Timber'), (1, 'overallcond'), (1, 'paveddrive_P'), (1, 'paveddrive_Y'), (1, 'roofmatl_Tar&Grv'), (1, 'roofmatl_WdShake'), (1, 'roofstyle_Gable'), (1, 'roofstyle_Gambrel'), (1, 'roofstyle_Mansard'), (1, 'saletype_ConLw'), (1, 'saletype_New'), (1, 'saletype_Oth'), (1, 'saletype_WD'), (1, 'utilities_NoSeWa'), (2, 'bsmtfintype2_No'), (3, 'bsmtfintype2_Unf'), (4, 'condition1_RRAe'), (5, 'bsmtcond_No'), (6, 'roofmatl_CompShg'), (7, 'bldgtype_2fmCon'), (8, 'neighborhood_BrkSide'), (9, 'age'), (10, 'yrsold'), (11, 'yearbuilt'), (12, 'foundation_Stone'), (13, 'bldgtype_TwnhsE'), (14, 'neighborhood_Veenker'), (15, 'saletype_ConLD'), (16, 'garagetype_No'), (17, 'lotshape_IR3'), (18, 'exterqual_Gd'), (19, 'exterqual_TA'), (20, 'condition2_Feetr'), (21, 'bsmtfintype1_Rec'), (22, 'condition1_PosA'), (23, 'exterior2nd_Brk Cmn'), (24, 'miscfeature_Shed'), (25, 'bsmtqual_No'), (26, 'roofstyle_Hip'), (27, 'housestyle_1Story'), (28, 'heating_OthW'), (29, 'bsmtfintype1_No'), (30, 'bsmtqual_Gd'), (31, 'kitchenabvgr'), (32, 'garagecond_No'), (33, 'exterior1st_Plywood'), (34, 'exterior2nd_MetalSd'), (35, 'poolqc_Fa'), (36, 'garagequal_No'), (37, 'exterior1st_MetalSd'), (38, 'exterior1st_BrkFace'), (39, 'poolqc_Gd'), (40, 'neighborhood_Sawyer'), (41, 'overallqual'), (42, 'foundation_PConc'), (43, 'bsmtexposure_No'), (44, 'garagefinish_Unf'), (45, 'remod_age'), (46, 'yearremodad')

```